Proposal: AI Back-Office Knowledge Agent with Continuous Learning

Project Overview

We propose the development of an intelligent AI agent that serves as a company's internal back-office assistant. Built on **Hermes Agent**, the solution leverages Hermes' native long-term memory capabilities to continuously learn from interactions while accessing the company's latest knowledge through integrated data sources.

Employees will interact with the agent directly in Slack, receiving accurate, context-aware answers based on company documentation, previous interactions, and organizational knowledge.

The solution is designed to be modular and scalable, allowing additional capabilities such as a centralized context database to be introduced as the organization's knowledge ecosystem grows.

Objectives

The AI agent will:

- Answer questions about company processes, policies, and internal documentation.
- Support employees with everyday back-office tasks.
- Learn continuously using Hermes Agent's native memory capabilities.
- Retrieve up-to-date information from company knowledge sources.
- Operate directly within Slack.
- Provide consistent, professional, and context-aware responses.
- Scale alongside the organization's growing documentation and knowledge.

Scope of Work

1. Hermes Agent Implementation

The project includes the deployment and configuration of Hermes Agent as the core AI platform.

This includes:

- Hermes Agent installation and configuration
- Tool configuration
- Native long-term memory configuration
- Reasoning workflow configuration
- Secure authentication
- Deployment configuration

Hermes Agent's built-in memory will be responsible for learning from conversations, remembering recurring context, and improving future interactions without requiring custom memory management.

---

2. Agent Persona Design

A well-defined persona is essential to ensure that employees experience the assistant as a reliable and consistent member of the organization rather than a generic chatbot.

The persona defines:

- Communication style
- Tone of voice
- Company terminology
- Response structure
- Decision boundaries
- Escalation behavior
- Handling of uncertainty
- Professional etiquette

By defining a clear persona, responses remain consistent across departments, users receive predictable support, and the assistant reflects the company's communication standards. This significantly improves user trust and adoption.

---

3. Knowledge Integration

The agent will retrieve information from the company's primary knowledge sources:

Notion

Notion will serve as the central knowledge base containing:

- Company policies
- SOPs
- Internal documentation
- FAQs
- Process documentation
- Team knowledge

OneDrive

Microsoft OneDrive will provide access to business documents such as:

- PDF manuals
- Word documents
- Excel files
- Presentations
- Technical documentation
- Shared project documents

Hermes Agent will retrieve relevant information from these sources before generating responses, ensuring employees receive accurate and current information.

---

4. Slack Integration

Employees will communicate with the assistant directly inside Slack.

Features include:

- Direct Messages
- Mention support
- Channel conversations
- Thread awareness
- Rich Slack formatting
- Optional slash commands
- Permission-aware responses

This allows employees to access company knowledge without leaving their daily workflow.

---

5. Optional Context Database Integration

As an optional enterprise enhancement, Hermes Agent can be connected to a dedicated context database such as **Honcho**, **OpenViking**, or **Supermemory**.

This component complements—not replaces—Hermes Agent's native memory.

Hermes Memory vs. Context Database

Hermes Agent Memory is responsible for:

- Learning from conversations
- Remembering interaction history
- Retaining user preferences
- Building long-term conversational context
- Improving future responses

Context Database is responsible for:

- Centralized organizational knowledge
- Semantic retrieval across multiple sources
- Synchronization with company documentation
- Persistent context shared across agents
- Efficient retrieval of large document collections

Instead of maintaining local Markdown (.md) knowledge files, the context database continuously indexes content from Notion and OneDrive, allowing the assistant to always retrieve the latest approved information.

Benefits over Local Markdown Files

Compared to local documentation, a context database provides:

- Automatic synchronization
- Centralized knowledge management
- Better semantic search
- Easier maintenance
- Support for large knowledge collections
- Shared context across multiple AI agents
- Reduced manual deployment effort
- Better scalability as documentation grows

This architecture separates **knowledge retrieval** from **conversation memory**, resulting in a more maintainable and enterprise-ready solution.

---

6. Security & Permissions

The solution will implement appropriate security controls, including:

- Role-based access control
- Secure authentication
- Encrypted communication
- Audit logging
- Access restrictions based on documentation permissions
- Compliance with company security policies

Deliverables

- Configured Hermes Agent
- Agent persona definition
- Slack integration
- Notion integration
- OneDrive integration
- Optional context database integration (Honcho, OpenViking, or Supermemory)
- Deployment configuration
- Administrator documentation
- User documentation

Suggested Technology Stack

Component	Technology
AI Agent	Hermes Agent
LLM	OpenAI (or customer-selected provider)

Component	Technology
Collaboration	Slack API
Knowledge Base	Notion API
Document Source	Microsoft Graph API (OneDrive)
Context Layer (Optional)	Honcho / OpenViking / Supermemory
Programming Language	Python
Containerization	Docker
Infrastructure	Cloud VM or Kubernetes (optional)

Expected Workflow

1. An employee asks a question in Slack.
2. Hermes Agent receives the request.
3. Relevant information is retrieved from Notion and OneDrive.
4. If configured, the context database provides additional semantic context across organizational knowledge.
5. Hermes Agent combines retrieved knowledge with its native long-term memory.
6. The agent generates an accurate, context-aware response.
7. Hermes Agent updates its conversational memory for future interactions.

Benefits

- Centralized access to company knowledge
- Faster employee onboarding
- Reduced workload for back-office teams
- Consistent and professional responses
- Native long-term learning through Hermes Agent
- Scalable knowledge architecture
- Seamless Slack experience
- Enterprise-ready knowledge management
- Modular architecture for future AI agents

Operational Cost Considerations

The following recurring costs should be considered during deployment and operation:

- LLM API usage (based on token consumption)
- Hermes Agent hosting and compute resources
- Slack application hosting
- Optional context database hosting (Honcho, OpenViking, or Supermemory)

- Storage requirements for indexed documents
- Synchronization jobs for Notion and OneDrive
- Embedding generation for newly added or updated documentation (if required by the selected context database)

Operational costs will primarily depend on the number of users, document volume, synchronization frequency, and the selected hosting model. Self-hosted context databases generally reduce subscription costs while increasing infrastructure management responsibilities, whereas managed services provide lower maintenance with predictable recurring fees.

Estimated Timeline

Phase	Duration
Architecture & Solution Design	1 week
Hermes Agent Setup & Persona Design	1 week
Notion & OneDrive Integration	1–2 weeks
Slack Integration	1 week
Optional Context Database Integration	1 week
Testing & Optimization	1 week
Documentation & Deployment	1 week

Estimated total duration: 6–8 weeks, depending on the complexity of the existing infrastructure and the selected optional components.

Conclusion

The proposed solution delivers a modern AI-powered back-office assistant that combines Hermes Agent's native conversational memory with direct access to organizational knowledge stored in Notion and OneDrive. Employees interact naturally through Slack while benefiting from accurate, context-aware responses that improve over time.

For organizations seeking greater scalability, an optional context database such as Honcho, OpenViking, or Supermemory provides a centralized knowledge layer that complements Hermes Agent's memory, enabling enterprise-grade knowledge retrieval without sacrificing maintainability. This architecture creates a future-proof foundation that can evolve into a broader AI platform supporting multiple agents and business workflows.